

301AA - Advanced Programming

Lecturer: **Andrea Corradini**

andrea@di.unipi.it

<http://pages.di.unipi.it/corradini/>

Course pages:

<http://pages.di.unipi.it/corradini/Didattica/AP-18/>

AP-2018-04: Runtime Systems and intro to JVM

Overview

- Runtime Systems
- The Java Runtime Environment
- The JVM as an abstract machine
- JVM Data Types
- JVM Runtime Data Areas
- Multithreading
- Per-thread Data Areas
- Dynamic Linking
- JIT compilation
- Method Area

Runtime system

- Every programming language defines an **execution model**
- A **runtime system** implements (part of) such execution model, providing support during the execution of corresponding programs
- **Runtime support** is needed both by *interpreted* and by *compiled* programs, even if typically less by the latter

Runtime system (2)

- The runtime system can be made of
 - Code in the executing program generated by the compiler
 - Code running in other threads/processes during program execution
 - Language libraries
 - Operating systems functionalities
 - The interpreter / virtual machine itself

Runtime Support needed for...

- Memory management
 - Stack management: Push/pop of activation records
 - Heap management: allocation, garbage collection
 - ➔ **Chapter 7 of "Dragon Book"**
- Input/Output
 - Interface to file system / network sockets / I/O devices
- Interaction with the **runtime environment**,
 - state values accessible during execution (eg. environment variables)
 - active entities like disk drives and people via keyboards.

Runtime Support needed for... (2)

- Parallel execution via threads/tasks/processes
- Dynamic type checking and dynamic binding
- Dynamic loading and linking of modules
- Debugging
- Code generation (for JIT compilation) and Optimization
- Verification and monitoring

Java Runtime Enviroment - JRE

- Includes all what is needed to run compiled Java programs
 - JVM – Java Virtual Machine
 - JCL – Java Class Library (Java API)
- We shall focus on the JVM as a real runtime system covering most of the functionalities just listed
- Reference documentation:
 - The Java™ Virtual Machine Specification, Java SE 8 Edition
<https://docs.oracle.com/javase/specs/jvms/se8/jvms8.pdf>
 - The Java Language Specification, Java SE 8 Edition
<https://docs.oracle.com/javase/specs/jls/se8/jls8.pdf>

New "short-term" releases of Java

➔ **Java 9**, released in September 2017

- Added module system
- <https://www.oracle.com/corporate/features/understanding-java-9-modules.html>

➔ **Java 10**, released in March 2018

- Type inference of local variables
- <https://developer.oracle.com/java/jdk-10-local-variable-type-inference>
- <https://medium.com/the-java-report/java-10-sneak-peek-local-variable-type-inference-var-3022016e1a2b>

➔ **Java 11 (LTS)**, to be released in September 2018

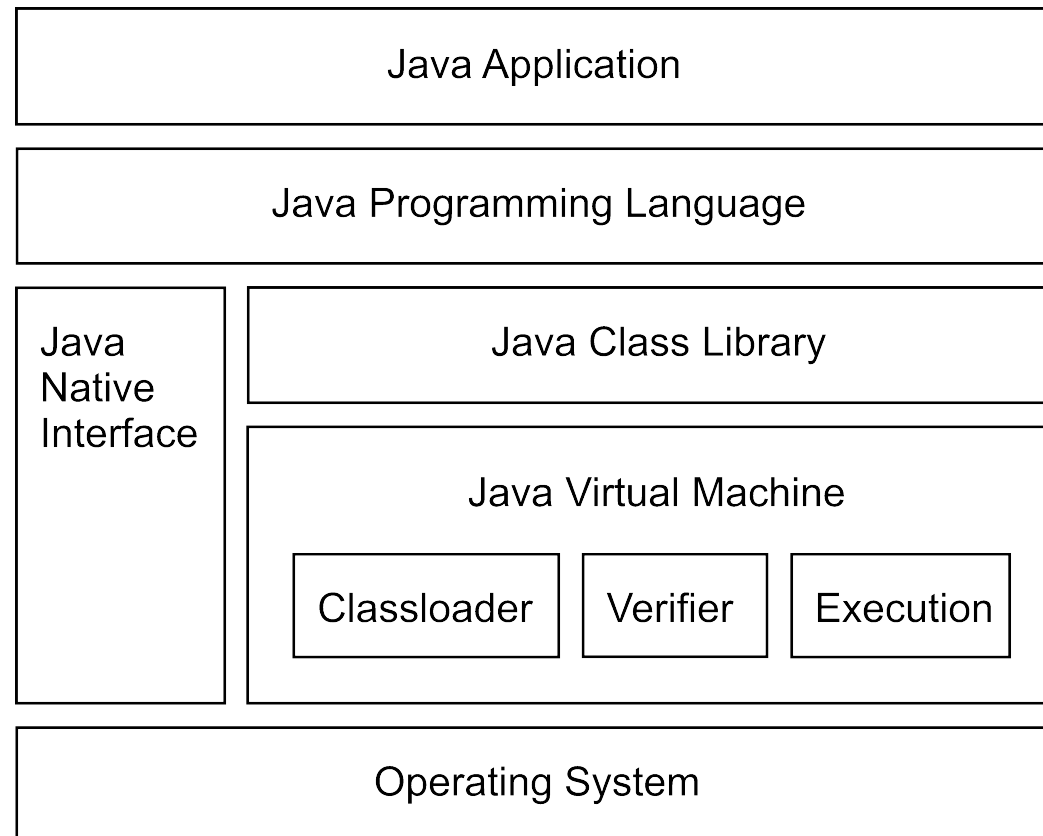
What is the JVM?

- The **JVM** is an **abstract** machine in the true sense of the word.
- The JVM specification does *not* give implementation details like memory layout of run-time data area, garbage-collection algorithm, internal optimization (can be dependent on target OS/platform, performance requirements, etc.)
- The JVM specification defines a machine independent “**class file format**” that all JVM implementations must support
- The JVM imposes **strong syntactic** and **structural constraints** on the code in a class file. Any language with functionality that can be expressed in terms of a valid class file can be hosted by the JVM

Execution model

- JVM is a *multi-threaded stack based machine*
- JVM instructions
 - implicitly take arguments from the top of the operand stack of the current frame
 - put their result on the top of the operand stack
- The operand stack is used to
 - pass arguments to methods
 - return a result from a method
 - store intermediate results while evaluating expressions
 - store local variables


Java Abstract Machine Hierarchy



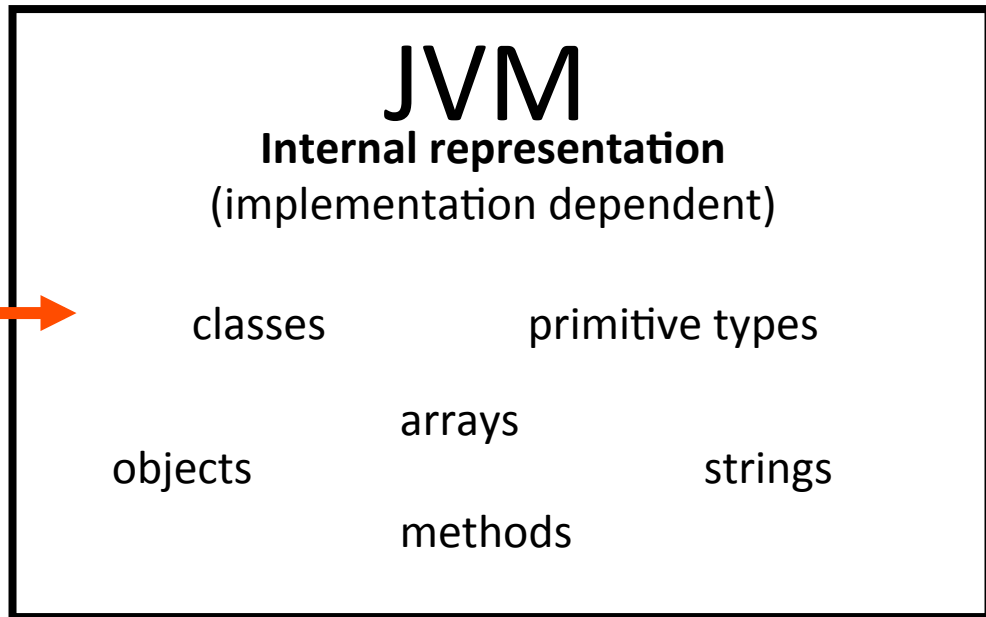
Class Files and Class File Format

External representation
(platform independent)

.class files



load



JVM Data Types

Primitive types:

- boolean: `boolean` (support only for arrays)
- numeric integral: `byte`, `short`, `int`, `long`, `char`
- numeric floating point: `float`, `double`
- internal, for exception handling: `returnAddress`

Reference types:

- class types
- array types
- interface types


Note:

- No type information on local variables at runtime
- Types of operands specified by opcodes (eg: `iadd`, `fadd`,)

Object Representation

- Left to the implementation
 - Including concrete value of `null`
- This adds extra level of indirection
 - need pointers to instance data and class data
 - make garbage collection easier
- Object representation must include
 - mutex lock
 - GC state (flags)

JVM Runtime Data Areas



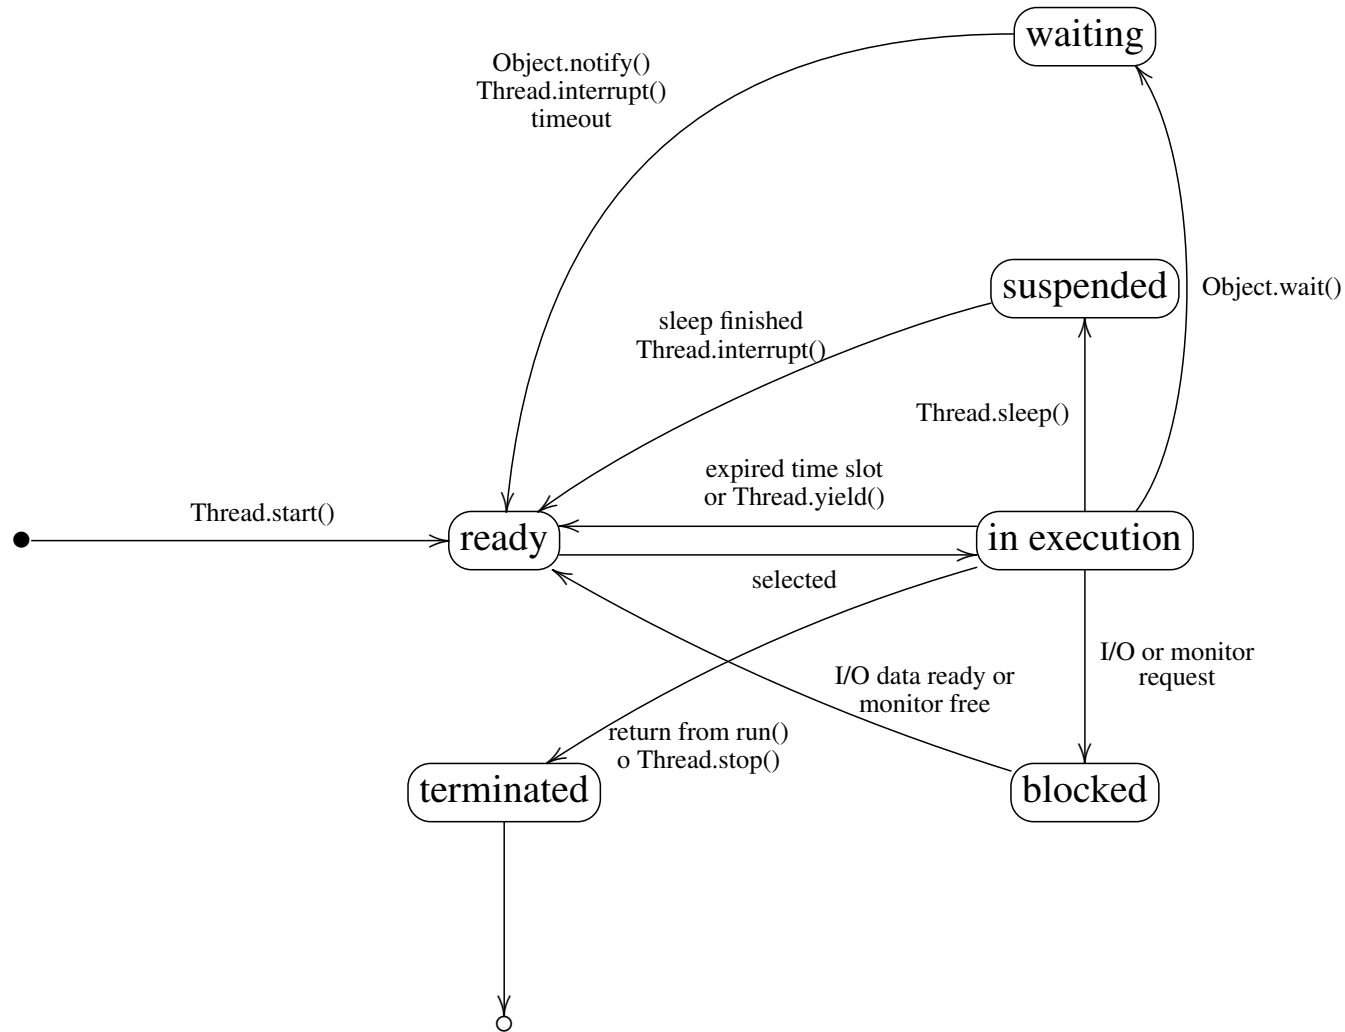
Threads

- JVM allows multiple threads per application, starting with `main`
- Created as instances of `Thread` invoking `start ()` (which invokes `run ()`)
- Several background (daemon) system threads for
 - Garbage collection, finalization
 - Signal dispatching
 - Compilation, etc.
- Threads can be supported by time-slicing and/or multiple processors

Threads (2)


- Threads have shared access to heap and persistent memory
- Complex specification of consistency model
 - volatiles
 - working memory vs. general store
 - non-atomic longs and doubles
- The *Java programming language memory model* prescribes the behaviour of multithreaded programs (JLS-8 Ch. 17)

Java Thread Life Cycle




Per Thread Data Areas

- **pc:** pointer to next instruction in *method area*
 - undefined if current method is *native*
- The **java stack:** a stack of *frames* (or *activation records*).
 - A new frame is created each time a method is invoked and it is destroyed when the method completes.
 - The JVMMS does not require that frames are allocated contiguously
- The **native stack:** is used for invocation of native functions, through the JNI (Java Native Interface)
 - When a native function is invoked, eg. a C function, execution continues using the native stack
 - Native functions can call back Java methods, which use the Java stack



Structure of frames



- **Local Variable Array** (32 bits) containing
 - Reference to `this` (if instance method)
 - Method parameters
 - Local variables
- **Operand Stack** to support evaluation of expressions and evaluation of the method
 - Most JVM bytecodes manipulate the stack
- Reference to **Constant Pool** of current class

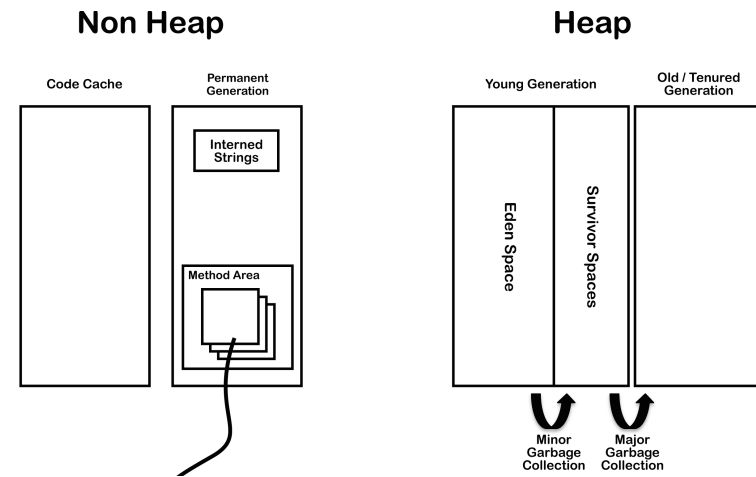
Dynamic Linking (1)

- The reference to the constant pool for the current class helps to support **dynamic linking**.
- In C/C++ typically multiple object files are linked together to produce an executable or dll.
 - During the linking phase symbolic references are replaced with an actual memory address relative to the final executable.
- In Java this linking phase is done **dynamically** at runtime.
- When a Java class is compiled, all references to variables and methods **are stored in the class's constant pool as symbolic references**.

Dynamic Linking (2)


- The JVM implementation can choose when to resolve symbolic references.
 - **Eager or static resolution:** when the class file is verified after being loaded
 - **Lazy or late resolution:** when the symbolic reference is used for the first time
- The JVM **has to behave** as if the resolution occurred when each reference is first used and throw any resolution errors at this point.
- **Binding** is the process of the entity (field, method or class) identified by the symbolic reference being replaced by a direct reference
- This only happens once because the symbolic reference *is completely replaced* in the constant pool
- If the symbolic reference refers to a class that has not yet been resolved then this class will be loaded.

Data Areas Shared by Threads: **Heap**



- Memory for objects and arrays; unlike C/C++ they are never allocated to stack
- Explicit deallocation not supported. Only by garbage collection.
- The HotSpot JVM includes four **Generational Garbage Collection Algorithms**

Data Areas Shared by Threads: **Non-Heap**



- Memory for objects which are never deallocated, needed for the JVM execution
 - Method area
 - Interned strings
 - Code cache for JIT

JIT compilation

- The Hotspot JVM (and other JVMs) profiles the code during interpretation, looking for “hot” areas of byte code that are executed regularly
- These parts are compiled to native code.
- Such code is then stored in the **code cache** in non-heap memory.

Method area

The memory where `class` files are loaded. For each class:

- **ClassLoader Reference**
- From the `class` file:
 - **Run Time Constant Pool**
 - **Field data**
 - **Method data**
 - **Method code**

Note: Method area is shared among thread. Access to it has to be **thread safe**.

Changes of method area when:

- A new class is loaded
- A symbolic link is resolved by dynamic linking

Class file structure

ClassFile {

u4	magic;	0xCAFEBAE
u2	minor_version;	Java Language Version
u2	major_version;	
u2	constant_pool_count;	Constant Pool
cp_info	constant_pool[constant_pool_count-1];	
u2	access_flags;	access modifiers and other info
u2	this_class;	References to Class and Superclass
u2	super_class;	
u2	interfaces_count;	References to Direct Interfaces
u2	interfaces[interfaces_count];	
u2	fields_count;	Static and Instance Variables
field_info	fields[fields_count];	
u2	methods_count;	Methods
method_info	methods[methods_count];	
u2	attributes_count;	Other Info on the Class
attribute_info	attributes[attributes_count];	

}

Field data in the Method Area

Per field:

- Name
- Type
- Modifiers
- Attributes

FieldType descriptors

<i>FieldType</i> term	Type	Interpretation
B	byte	signed byte
C	char	Unicode character code point in the Basic Multilingual Plane, encoded with UTF-16
D	double	double-precision floating-point value
F	float	single-precision floating-point value
I	int	integer
J	long	long integer
L <i>ClassName</i> ;	reference	an instance of class <i>ClassName</i>
S	short	signed short
Z	boolean	true or false
[reference	one array dimension

Method data

Per method:

- Name
- Return Type
- Parameter Types (in order)
- Modifiers
- Attributes

A *method descriptor* contains

- a sequence of zero or more *parameter descriptors* in brackets
- a *return descriptor* or V for *void descriptor*

Example: The descriptor of

```
Object m(int i, double d, Thread t) {...}
```

is:

```
(IDLjava/lang/Thread;)Ljava/lang/Object;
```