301AA - Advanced Programming

Lecturer: **Andrea Corradini**
andrea@di.unipi.it
http://pages.di.unipi.it/corradini/

Course pages:
http://pages.di.unipi.it/corradini/Didattica/AP-18/

AP-2018-23: *Streams in Java 8*

Java 8: language extensions

Java 8 is the biggest change to Java since the inception of the language. Main new features:

- Lambda expressions
  - Method references
  - Default methods in interfaces
  - Improved type inference

- Stream API

A big challenge was to introduce lambdas without requiring recompilation of existing binaries

Streams in Java 8

The **java.util.stream** package provides utilities to support functional-style operations on streams of values. **Streams** differ from collections in several ways:

- **No storage**. A stream is not a data structure that stores elements; instead, it conveys elements from a *source* (a data structure, an array, a generator function, an I/O channel,...) through a *pipeline* of computational operations.

- **Functional in nature**. An operation on a stream produces a result, but does not modify its source.

Streams in Java 8 (cont'd)

- **Laziness-seeking**. Many stream operations, can be implemented lazily, exposing opportunities for optimization. Stream operations are divided into **intermediate** (stream-producing) operations and **terminal** (value- or side-effect-producing) operations. ***Intermediate operations are always lazy.***

- **Possibly unbounded**. Collections have a finite size, streams need not. Short-circuiting operations such as ***limit(n)*** or ***findFirst()*** can allow computations on infinite streams to complete in finite time.

- **Consumable**. The elements of a stream are only visited once during the life of a stream. Like an *Iterator*, a new stream must be generated to revisit the same elements of the source.

Pipelines

- A typical pipeline contains
  - A **source**, producing (by need) the elements of the stream
  - Zero or more **intermediate operations**, producing streams
  - A **terminal operation**, producing side-effects or non-stream values
- Example of typical pattern: filter / map / reduce

```
double average = listing   // collection of Person
    .stream()              // stream wrapper over a collection
    .filter(p -> p.getGender() == Person.Sex.MALE) // filter
    .mapToInt(Person::getAge)     // extracts stream of ages
    .average()              // computes average (reduce/fold)
    .getAsDouble();        // extracts result from OptionalDouble
```

Anatomy of the Stream Pipeline

- A Stream is processed through a pipeline of operations
- A Stream starts with a source
- Intermediate methods are performed on the Stream elements. These methods produce Streams and are not processed until the terminal method is called.
- The Stream is considered consumed when a terminal operation is invoked. No other operation can be performed on the Stream elements afterwards
- A Stream pipeline may contain some short-circuit methods (which could be intermediate or terminal methods) that cause the earlier intermediate methods to be processed only until the short-circuit method can be evaluated.

Stream sources

Streams can be obtained in a number of ways:

- From a **Collection** via the **stream()** and **parallelStream()** methods;

- From an array via **Arrays.stream(Object[])**;

- From static factory methods on the stream classes, such as **Stream.of(Object[])**, **IntStream.range(int, int)** or **Stream.iterate(Object, UnaryOperator)**;

- The lines of a file can be obtained from **BufferedReader.lines()**;

- Streams of file paths can be obtained from methods in **Files**;

- Streams of random numbers can be obtained from **Random.ints()**;

- Generators, like **generate** or **iterate**;

- Numerous other methods in the JDK…

Intermediate Operations

- An intermediate operation keeps a stream open for further operations. Intermediate operations are lazy.
- Several intermediate operations have arguments of *functional interfaces*, thus *lambdas* can be used

```
Stream<T> filter(Predicate<? super T> predicate)  // filter

IntStream mapToInt(ToIntFunction<? super T> mapper) // map f:T -> int

<R> Stream<R> map(Function<? super T,? extends R> mapper) // map f:T->R

Stream<T> peek(Consumer<? super T> action) //performs action on elements

Stream<T> distinct() // remove duplicates – stateful

Stream<T> sorted() // sort elements of the stream – stateful

Stream<T> limit(long maxSize) // truncate

Stream<T> skip(long n) // skips first n elements
```

Using peek…

- **peek** does not affect the stream
- A typical use is for debugging

```
IntStream.of(1, 2, 3, 4)
        .filter(e -> e > 2)
        .peek(e -> System.out.println("Filtered value: " + e))
        .map(e -> e * e)
        .peek(e -> System.out.println("Mapped value: " + e))
        .sum();
```

Terminal Operations

- A **terminal operation** must be the final operation on a stream. Once a terminal operation is invoked, the stream is consumed and is no longer usable.
- Typical: collect values in a data structure, reduce to a value, print or other side effects.

```
void forEach(Consumer<? super T> action)

Object[] toArray()

T reduce(T identity, BinaryOperator<T> accumulator)  // fold

Optional<T> reduce(BinaryOperator<T> accumulator) // fold

Optional<T> min(Comparator<? super T> comparator)

boolean allMatch(Predicate<? super T> predicate) // short-circuiting

boolean anyMatch(Predicate<? super T> predicate) // short-circuiting

Optional<T> findAny() // short-circuiting
```

Types of Streams

- Streams only for reference types, int, long and double
  - Minor primitive types are missing

```
"Hello world!".chars()
    .forEach(System.out::print);

// prints
72101108108111321119111114108108033


// fixing it:
"Hello world!".chars()
    .forEach(x -> System.out.print((char) x));
```

From Reduce to Collect: Mutable Reduction

- Suppose we want to concatenate a stream of strings.
- The following works but is highly inefficient (it builds one new string for each element):

```
String concatenated = listOfStrings
                          .stream()
                          .reduce("", String::concat)
```

- Better to "accumulate" the elements in a mutable object (a StringBuilder, a collection, …)
- The *mutable reduction operation* is called **collect()**. It requires three functions:
  - a supplier function to construct new instances of the result container,
  - an accumulator function to incorporate an input element into a result container,
  - a combining function to merge the contents of one result container into another.

```
<R> R collect( Supplier<R> supplier,
               BiConsumer<R, ? super T> accumulator,
               BiConsumer<R, R> combiner);
```

Mutable reductions: examples

- Collecting the String representations of the elements of a stream into an ArrayList:

```
// no streams
ArrayList<String> strings = new ArrayList<>();
for (T element : stream) {
    strings.add(element.toString());
}
```

```
// with streams and lambdas
ArrayList<String> strings =
    stream.collect(() -> new ArrayList<>(), //Supplier
    (c, e) -> c.add(e.toString()),    // Accumulator
    (c1, c2) -> c1.addAll(c2));       //Combining
```

```
// with streams and method references
ArrayList<String> strings = stream.map(Object::toString)
    .collect(ArrayList::new, ArrayList::add, ArrayList::addAll);
```

Mutable reductions: Collectors

- Method collect can also be invoked with a Collector argument:

```
<R,A> R collect(Collector<? super T,A,R> collector)
```

- A Collector encapsulates the functions used as arguments to collect(Supplier, BiConsumer, BiConsumer), allowing for reuse of collection strategies and composition of collect operations.

```java
// The following will accumulate strings into an ArrayList:

List<String> asList = stringStream.collect(Collectors.toList());

// The following will classify Person objects by city:

Map<String, List<Person>> peopleByCity =
personStream.collect(Collectors.groupingBy(Person::getCity));
```

Infinite Streams

- Streams wrapping collections are finite
- Infinite streams can be generated with:
  - iterate
  - generate

```
static <T> Stream<T> iterate(T seed, UnaryOperator<T> f)

// Example: summing first 10 elements of an infinite stream
int sum = Stream.iterate(0,x -> x+1).limit(10).reduce(0,(x,s) -> x+s);

static <T> Stream<T> generate(Supplier<T> s)

// Example: printing 10 random mumbers
Stream.generate(Math::random).limit(10).forEach(System.out::println);
```

Parallelism

- Streams facilitate parallel execution
- Stream operations can execute either in serial (default) or in parallel

```
double average = persons    //average age of all male
    .parallelStream()       // members in parallel
    .filter(p -> p.getGender() == Person.Sex.MALE)
    .mapToInt(Person::getAge)
    .average()
    .getAsDouble();
```

- The runtime support takes care of using multithreading for parallel execution, in a transparent way
- If operations don't have side-effects, **thread-safety is guaranteed** even if non-thread-safe collections are used (e.g.: ArrayList)

Parallelism (2)

- Concurrent mutable reduction supported for parallel streams
  - Suitable methods of Collector
- Order of processing stream elements depends on serial/parallel execution and intermediate operations

```
Integer[] intArray = {1, 2, 3, 4, 5, 6, 7, 8 };
List<Integer> listOfIntegers = new ArrayList<>(Arrays.asList(intArray));
    listOfIntegers .stream()
                    .forEach(e -> System.out.print(e + " "));
        // prints: 1 2 3 4 5 6 7 8
    listOfIntegers .parallelStream()
                    .forEach(e -> System.out.print(e + " "));
        // may print: 3 4 1 6 2 5 7 8
```

A simple parallel stream example

Slide by Josh Bloch

- Consider this for-loop (.96 s runtime; dual-core laptop)

```
long sum = 0;
for (long j = 0; j < Integer.MAX_VALUE; j++) sum += j;
```

- Equivalent stream computation (1.5 s)

```
long sum = LongStream.range(0, Integer.MAX_VALUE).sum();
```

- Equivalent parallel computation (.77 s)

```
long sum = LongStream.range(0,Integer.MAX_VALUE)
          .parallel().sum();
```

- Fastest handcrafted parallel code I could write (.48 s)

  – You don't want to see the code. It took hours.

When to use a parallel stream – loosely speaking

- When operations are independent, and
- Either or both:
  - Operations are computationally expensive
  - Operations are applied to many elements of efficiently splittable data structures
- **Always measure before and after parallelizing!**
  - Jackson's third law of optimization

SplitIterator: Streams from collections

- A stream wrapping a collection uses a **Splititerator** over the collection

- This is the parallel analogue of an **Iterator**: it describes a (possibly infinite) collection of elements with support for
  - sequentially advancing,
  - *applying an action* to the next or to all remaining elements
  - splitting off some portion of the input into another spliterator which can be processed in parallel.

- At the lowest level, all streams are driven by a spliterator.

When to use a parallel stream – in detail

Slide by Josh Bloch

- Consider `s.parallelStream().operation(f)` if
  - `f`, the per-element function, is independent
    - i.e., computation for each element doesn't rely on or impact any other
  - `s`, the source collection, is efficiently splittable
    - Most collections, and `java.util.SplittableRandom`
    - NOT most I/O-based sources
  - Total time to execute sequential version roughly > 100µs
    - "Multiply N (number of elements) by Q (cost per element of `f`), guestimating Q as the number of operations or lines of code, and then checking that N*Q is at least 10,000.

Critical issues

- Non-interference
  - Behavioural parameters (like lambdas) of stream operations should not affect the source (***non-interfering behaviou**r*)
  - Risk of **ConcurrentModificationExceptions**, even if in single thread
- Stateless behaviours
  - Statless behaviour for intermediate operations is encouraged, as it facilitates parallelism, and functional style, thus maintenance
- Parallelism and thread safety
  - For parallel streams with side-effects, ensuring thread safety is the programmers' responsibility

Inteference: an example

```java
try {
    List<String> listOfStrings =
        new ArrayList<>(Arrays.asList("one", "two"));

    String concatenatedString = listOfStrings
        .stream()
// Don't do this! Interference occurs here.
        .peek(s -> listOfStrings.add("three"))
        .reduce((a, b) -> a + " " + b)
        .get();
    System.out.println("Concatenated string: " + concatenatedString);
} catch (Exception e) {
    System.out.println("Exception caught: " + e.toString());
}
```

MONADS IN JAVA….

Monads in Java: Optional and Stream

```java
public static <T> Optional<T> of(T value)
// Returns an Optional with the specified present non-null value.

<U> Optional<U> flatMap(Function<? super T,Optional<U>> mapper)
/* If a value is present, apply the provided Optional-bearing mapping
function to it, return that result, otherwise return an empty
Optional. */
```

```java
static <T> Stream<T> of(T t)
// Returns a sequential Stream containing a single element.

<R> Stream<R> flatMap(
        Function<? super T,? extends Stream<? extends R>> mapper)
/* Returns a stream consisting of the results of replacing each element
of this stream with the contents of a mapped stream produced by applying
the provided mapping function to each element.  */
```

Functional programming and monads in Java

- About the way monads entered the Java landscape I suggest reading the slides on **Monadic Java** by Mario Fusco.

- More on functional programming in Java in the book **Java 8 in action**